

# **NGINX Cheat Sheet , From Basic to Advanced**

A compact reference of common and advanced NGINX configuration snippets, with short explanations and recommended options. Use as a quick reference during workshops or when building configs.

## Table of Contents

- Beginner-friendly walkthrough
- What the nginx commands and common directives mean (CLI reference)
- Minimal nginx.conf skeleton
- Basic server block (static site)
- Running NGINX in Docker (optional)
- Reverse proxy
- Upstream / Load balancing
- Proxy tuning
- Caching
- Rate limiting
- Advanced topics (Lua, map, sticky sessions)
- Troubleshooting & debugging

## **Quick commands**

- See the CLI section below: "What the nginx commands and common directives mean" for detailed explanation and examples (test/reload and Docker commands).

## **What the nginx commands and common directives mean**

This short reference explains the CLI commands and the most-used configuration directives in this cheat sheet. It helps you understand *what* each command does, *why* you'd run it, and *what to expect*.

## CLI commands

- `nginx -t`
  - Purpose: Test the active NGINX configuration for syntax errors and validate includes.
  - What it does: Parses the configuration files (`nginx.conf` and included files) and reports syntax correctness and any missing files or invalid directives.
  - When to use: Always run after editing any config and before reloading/restarting NGINX.
  - Expected output: `nginx: the configuration file /etc/nginx/nginx.conf test is successful` on success; errors list file and line on failure.
- `nginx -s reload`
  - Purpose: Gracefully reload NGINX workers with new configuration without dropping connections.
  - What it does: Sends a signal to the master process to reload the configuration; master spawns new worker processes with the new config and gracefully shuts down old workers.
  - When to use: After a successful `nginx -t`, to apply changes without downtime.
  - Note: On systemd systems you can use `systemctl reload nginx` which issues the same graceful reload.

- `docker exec <container> nginx -t` / `docker exec <container> nginx -s reload`
  - Purpose: Run the same `nginx` CLI commands inside a running container.
  - What it does: Executes the command inside the container namespace; useful when NGINX runs in Docker and you don't have a native package-installed binary.
  - When to use: Test/reload containerized NGINX after replacing mounted configs or updating image files.

## Common NGINX directives (what they do)

- `server { ... }`
  - Purpose: Define a virtual host (listener) for a combination of IP/port and `server_name`.
  - Goal: Route requests for a domain or port to the relevant `location` blocks and behavior.
- `listen 80; / listen 443;`
  - Purpose: Specify which port and (optionally) IP address the server block should accept connections on.
  - Goal: Bind NGINX to that port. Port 80 for HTTP, 443 for HTTPS (TLS) traditionally.
- `server_name example.com;`
  - Purpose: Match the `Host` header to select which server block handles the request.
- `location / { ... }`
  - Purpose: Match request URIs and apply rules (proxying, serving static files, rewrites).
  - Goal: Map URL paths to backends, static content, or special handlers.
- `root /path/to/html; and index index.html;`
  - Purpose: Serve static files from the filesystem under `root` and default file names using `index`.

- `try_files $uri $uri/ =404;`
  - Purpose: Try to serve the requested file or fall back to the next expression; `=404` returns 404 if missing.
  - Goal: Provide clean static file serving with fallback behavior.
- `proxy_pass http://backend;`
  - Purpose: Forward the client request to an upstream/backend server or upstream group defined by `upstream`.
  - What to watch: Relative vs absolute URIs can change how the request path is forwarded — see official docs for subtle behavior.
- `upstream backend { server ... }`
  - Purpose: Define a named pool of backend servers and parameters (weights, fail policies).
  - Goal: Use the pool in `proxy_pass http://backend;` for load balancing.
- `proxy_set_header ...`
  - Purpose: Add or replace request headers when proxying to preserve client info (Host, X-Real-IP, X-Forwarded-For, etc.).
  - Goal: Ensure backend apps can see the original client address and host.

- `proxy_http_version 1.1;` and `proxy_set_header Connection "";`
  - Purpose: Use HTTP/1.1 for upstream connections (required for keepalive and WebSocket proxying) and clear `Connection` header to avoid interfering with connection management.
- `proxy_buffering off;`
  - Purpose: Disable proxy response buffering so data streams directly from upstream to client.
  - When to use: Real-time streams, server-sent events, or WebSocket-like flows.
- `limit_req_zone` and `limit_req` / `limit_conn_zone` and `limit_conn`
  - Purpose: Define rate limits and concurrent connection limits by a key (IP, cookie, etc.).
  - Goal: Protect upstreams from abusive traffic and reduce the blast radius of floods.
- `proxy_cache_path` and `proxy_cache` / `proxy_cache_valid`
  - Purpose: Configure on-disk cache for proxied responses to speed up repeated requests.
  - Goal: Reduce backend load and improve response times for cacheable content.

- `add_header NAME "value" always;`
  - Purpose: Add HTTP response headers (security headers, caching headers) to responses.
  - Note: `always` ensures the header is present for error responses too.
- `log_format` and `access_log` and `error_log`
  - Purpose: Define structured log formats and where to write logs. Essential for troubleshooting and metrics.

# Beginner-friendly walkthrough (line-by-line explanations)

This section explains the config snippets in plain language for someone who is new to NGINX. Each item below references a short snippet from the cheat sheet and explains the most important lines and why they are there.

## 1) Minimal `nginx.conf` skeleton

```
user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log warn;
pid /var/run/nginx.pid;

events { worker_connections 1024; }

http {
    include mime.types;
    default_type application/octet-stream;
    sendfile on;
    keepalive_timeout 65;
    include /etc/nginx/conf.d/*.conf;
}
```

- `user nginx;` — the OS user NGINX workers run as for permissions and security. On containers this often maps to a different UID.
- `worker_processes auto;` — how many worker processes NGINX should start. `auto` uses CPU cores automatically.
- `error_log /var/log/nginx/error.log warn;` — write NGINX internal errors to a file (useful for debugging).
- `events { worker_connections 1024; }` — controls how many connections each worker can handle. Multiply by `worker_processes` for total concurrency (rough estimate).

- `http { ... }` — top-level block for HTTP-specific settings. `include` brings in separate site files so you can keep configs organized.
- `sendfile on;` — improves static file performance by using optimized OS calls.

Why this matters: Start with this skeleton. You will add `server {}` blocks under `conf.d/` or `sites-enabled/` for each website or virtual host.

## 2) Basic server block (static site)

```
server {
    listen 80;
    server_name example.com www.example.com;
    root /var/www/example.com/html;
    index index.html index.htm;
    location / { try_files $uri $uri/ =404; }
}
```

- `listen 80;` — accept HTTP requests on port 80.
- `server_name` — matches the Host header in incoming requests; multiple names can be listed.
- `root` — folder on disk where NGINX looks for files to serve.
- `index` — filenames to try when a client requests a directory (`/` → `index.html`).
- `location / { ... }` — rules for requests to the root path. `try_files` checks for the file and returns 404 if not found.

Beginner note: Create the `index.html` file in the `root` folder and then visit your server IP or mapped port to see it served by NGINX.

# Running NGINX in Docker

This section shows common ways to run the official NGINX image, explains the Docker CLI options used, and gives a lightweight `docker-compose` example. Start here if you plan to run NGINX inside a container.

## 1. Run official image (quick test)

```
docker run --name nginx-test -d -p 8080:80 nginx:stable
```

Explanation of flags:

- `docker run` : start a new container from an image.
- `--name nginx-test` : assign a container name to make it easier to reference.
- `-d` : detach; run the container in the background.
- `-p 8080:80` : map host port 8080 to container port 80 (host:container).
- `nginx:stable` : the image (official NGINX stable release). You can use `nginx:alpine` for a smaller image.

Now visit `http://localhost:8080` to see the default NGINX welcome page.

## 2. Run with custom config and content (bind mounts)

Prepare a local folder structure (example):

```
./nginx/
  |- conf.d/
  |  \_ default.conf
  |- html/
  |  \_ index.html
```

Run the container with mount points:

```
docker run -d --name nginx-site \
-p 80:80 \
-v $(pwd)/nginx/conf.d:/etc/nginx/conf.d:ro \
-v $(pwd)/nginx/html:/usr/share/nginx/html:ro \
nginx:stable
```

Notes:

- `-v <host>:<container>:ro` mounts a host directory into the container as read-only.
- Mounting `conf.d` and the site content lets you edit locally and immediately test by reloading the containerized NGINX.

### 3. Test and reload config inside the container

Test config:

```
docker exec nginx-site nginx -t
```

Reload gracefully:

```
docker exec nginx-site nginx -s reload
```

If you prefer a single command to replace files and reload, use `docker cp` to copy files into the container and then `nginx -s reload`.

#### 4. Build a custom image (Dockerfile)

If you need modules or pre-installed assets, build a custom image:

```
FROM nginx:stable
COPY conf.d/ /etc/nginx/conf.d/
COPY html/ /usr/share/nginx/html/

# Optional: copy custom mime types or other snippets
COPY snippets/ /etc/nginx/snippets/

EXPOSE 80
```

Build and run:

```
docker build -t my-nginx:local .
docker run -d --name my-nginx -p 80:80 my-nginx:local
```

## 5. Docker Compose example (recommended for multi-container setups)

`docker-compose.yml` :

```
version: '3.8'
services:
  nginx:
    image: nginx:stable
    container_name: nginx-compose
    ports:
      - "80:80"
    volumes:
      - ./nginx/conf.d:/etc/nginx/conf.d:ro
      - ./nginx/html:/usr/share/nginx/html:ro
    restart: unless-stopped
```

Commands:

- `docker-compose up -d` — start services in background.
- `docker-compose logs -f nginx` — follow NGINX logs.
- `docker-compose exec nginx nginx -t` — test config inside the service.

## 6. Permissions and SELinux/AppArmor

When bind-mounting, ensure the container user (usually `nginx` with UID 101 in official images) can read mounted files. On SELinux-enabled hosts, add `:z` or `:Z` to volume flags (e.g., `-v ./nginx/conf.d:/etc/nginx/conf.d:ro,z`).

## 7. Useful Docker commands for maintenance

- `docker ps` — list running containers.
- `docker logs -f nginx-site` — stream logs.
- `docker exec -it nginx-site /bin/bash` (or `/bin/sh` on alpine) — open a shell.
- `docker rm -f nginx-site` — stop and remove a container.

### 3) Reverse proxy (forwarding requests to an app)

```
location / {  
    proxy_pass http://127.0.0.1:3000;  
    proxy_set_header Host $host;  
    proxy_set_header X-Real-IP $remote_addr;  
}
```

- `proxy_pass` — tells NGINX to forward matching requests to another server (a backend application). Here the backend runs on the same machine at port 3000.
- `proxy_set_header Host $host;` — forwards the original Host header so the backend knows the requested domain.
- `proxy_set_header X-Real-IP $remote_addr;` — forwards the client's IP address so the backend can log or apply rate limits.

Beginner tip: Use this when your app listens on a port (e.g., a Node/Express app on 3000). NGINX will accept public requests and forward them to the app.

## 4) Upstream block (load balancing)

```
upstream backend {  
    server 10.0.0.11:8000 weight=3;  
    server 10.0.0.12:8000;  
}
```

- `upstream` — a named group of backend servers; use `proxy_pass http://backend;` to use them.
- `weight=3` — this server will receive roughly 3x the requests of a server without weight.

Beginner scenario: Start with two app instances and add them here. NGINX will distribute requests (round-robin by default).

## 5) Proxy tuning (timeouts and buffers)

See the dedicated "Proxy tuning (common recommended options)" section below for the full recommended example and explanations. Start with modest timeouts (connect 5s, read/send 30-60s) and only change buffer sizes when you observe header or response fragmentation.

## 6) Caching (`proxy_cache`)

```
proxy_cache_path /var/cache/nginx levels=1:2 keys_zone=STATIC:10m max_size=1g inactive=60m;
proxy_cache STATIC;
proxy_cache_valid 200 302 10m;
```

- `proxy_cache_path` — defines where cached files go on disk and how large the cache is.
- `proxy_cache` — enables caching in a location using the named zone.
- `proxy_cache_valid` — how long to keep cached responses for certain status codes.

Beginner benefit: Caching speeds up responses and reduces load on your application servers for static or rarely-changing content.

## 7) Rate limiting (protect your app)

```
limit_req_zone $binary_remote_addr zone=one:10m rate=10r/s;
limit_req zone=one burst=20 nodelay;
```

- `limit_req_zone` — creates a shared memory zone keyed by IP to count requests.
- `rate=10r/s` — allow roughly 10 requests per second per IP.
- `burst=20` — allow short spikes up to 20 requests.

Beginner warning: Set conservative limits and test; overly aggressive limits will block real users.

## 8) WebSocket proxying

```
proxy_http_version 1.1;
proxy_set_header Upgrade $http_upgrade;
proxy_set_header Connection "upgrade";
```

- WebSockets require HTTP/1.1 and the `Upgrade` header to be forwarded to the backend. Without these lines, WebSocket connections will fail.

## 9) Gzip compression

```
gzip on;
gzip_types text/plain text/css application/json application/javascript;
```

- `gzip on;` — compresses responses to save bandwidth. Useful for text-based responses.

Beginner tip: Test with `curl -I -H "Accept-Encoding: gzip" http://.../` and check `Content-Encoding: gzip` in the response.

## 10) Logging and troubleshooting

See the consolidated "Troubleshooting tips" section later in this file for validation, log locations, and practical `curl` checks to debug virtual hosts and upstreams.

## 11) Error pages and maintenance mode

`error_page 503 /maintenance.html;` lets you present a friendly page when you intentionally return a 503 (maintenance).

Beginner workflow: edit config -> run `nginx -t` -> `nginx -s reload` -> test URLs. If NGINX won't start, check `error.log` for the first reported error and fix that file/line.

## Minimal nginx.conf skeleton

```
user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log warn;
pid /var/run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    include      mime.types;
    default_type application/octet-stream;

    sendfile      on;
    tcp_nopush    on;
    tcp_nodelay   on;
    keepalive_timeout 65;

    include /etc/nginx/conf.d/*.conf;
    include /etc/nginx/sites-enabled/*;
}
```

Notes: `worker_processes auto` lets NGINX pick the number based on CPU cores. `sendfile` improves static file throughput.

## Basic server block (static site)

```
server {
    listen 80;
    server_name example.com www.example.com;

    root /var/www/example.com/html;
    index index.html index.htm;

    location / {
        try_files $uri $uri/ =404;
    }

    access_log /var/log/nginx/example.access.log;
}
```

## Reverse proxy (simple)

```
server {
    listen 80;
    server_name api.example.com;

    location / {
        proxy_pass http://127.0.0.1:3000;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
    }
}
```

Notes: forward standard headers so backend apps can see original client info.

## Upstream block — load balancing

```
upstream backend {
    server 10.0.0.11:8000 weight=3;
    server 10.0.0.12:8000;
    server 10.0.0.13:8000 max_fails=3 fail_timeout=30s;
    # keepalive 16; # enable persistent connections to backends
}

server {
    listen 80;
    server_name app.example.com;

    location / {
        proxy_pass http://backend;
        proxy_set_header Host $host;
        proxy_http_version 1.1;
        proxy_set_header Connection "";
    }
}
```

### Notes:

- `weight` changes traffic share.
- `max_fails` and `fail_timeout` form a simple failure policy.
- `keepalive` in upstream enables connection reuse (good for many short requests).

## Load balancing algorithms

- Round robin (default): list servers in `upstream {}`
- Least connections: `upstream backend { least_conn; server ... }`
- IP hash (session stickiness by client IP): `upstream backend { ip_hash; server ... }`
- Hash: `hash $request_uri consistent;` (requires the `hash` directive)

Note: Sticky sessions require third-party modules for cookie-based stickiness, or `ip_hash` for IP-based affinity.

## Health checks

- NGINX Open Source: no native active health checks; rely on `maxfails` / `fail_timeout` passive checks or use third-party modules (`nginx_upstream_check_module`) or external tools.
- NGINX Plus: has active health checks (`health_check` directive) and advanced features.

Passive example (in upstream): shown above (`maxfails` , `fail_timeout` ).

## Proxy tuning (common recommended options)

```
proxy_set_header Host $host;
proxy_set_header X-Real-IP $remote_addr;
proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
proxy_set_header X-Forwarded-Proto $scheme;
proxy_connect_timeout 5s;
proxy_send_timeout 60s;
proxy_read_timeout 60s;
proxy_buffer_size 16k;
proxy_buffers 4 32k;
proxy_busy_buffers_size 64k;
proxy_temp_file_write_size 64k;
```

Notes: adjust buffers/timeouts for your workload; too small causes fragmentation, too large wastes RAM.

## Caching (proxy cache)

```
http {
    proxy_cache_path /var/cache/nginx levels=1:2 keys_zone=STATIC:10m max_size=1g inactive=60m use_temp_path=off;

    server {
        location / {
            proxy_cache STATIC;
            proxy_cache_valid 200 302 10m;
            proxy_cache_valid 404 1m;
            proxy_cache_bypass $http_cache_control; # bypass when client asks
            proxy_pass http://backend;
        }
    }
}
```

Notes: tune `keys_zone` and `max_size` for available RAM/disk. `use_temp_path=off` reduces disk IO on some setups.

## Rate limiting (basic)

```
http {
    limit_req_zone $binary_remote_addr zone=one:10m rate=10r/s;

    server {
        location /api/ {
            limit_req zone=one burst=20 nodelay;
            proxy_pass http://backend;
        }
    }
}
```

Notes: `limit_req` protects against request floods (r/s). `burst` allows short spikes.

Connection limiting:

```
limit_conn_zone $binary_remote_addr zone=addr:10m;

location / {
    limit_conn addr 10;
}
```

## Gzip compression

```
http {
  gzip on;
  gzip_types text/plain text/css application/json application/javascript text/xml application/xml application/xml+rss text/javascript;
  gzip_min_length 256;
  gzip_proxied any;
}
```

## WebSocket proxying

```
location /ws/ {
    proxy_pass http://websocket_backend;
    proxy_http_version 1.1;
    proxy_set_header Upgrade $http_upgrade;
    proxy_set_header Connection "upgrade";
    proxy_set_header Host $host;
}
```

Notes: must use `proxy_http_version 1.1` and set `Upgrade` / `Connection` headers.

## Security headers (example)

```
add_header X-Frame-Options "SAMEORIGIN" always;
add_header X-Content-Type-Options "nosniff" always;
add_header X-XSS-Protection "1; mode=block" always;
add_header Referrer-Policy "no-referrer-when-downgrade" always;
add_header Content-Security-Policy "default-src 'self'; script-src 'self' 'unsafe-inline' 'unsafe-eval' https:" always;
```

Notes: CSP is powerful but can break pages if too strict.

## Client body and upload limits

```
client_max_body_size 50M;  
client_body_timeout 12s;
```

## Timeouts and keepalives

```
keepalive_timeout 65;
send_timeout 30s;
client_header_timeout 10s;
client_body_timeout 10s;
```

## Buffering control (for streaming / large responses)

- `proxy_buffering off;` disables buffering (useful for real-time streaming)
- `proxy_buffers`, `proxy_buffer_size`, and `proxy_busy_buffers_size` control memory used for buffered responses

Example to disable buffering for a location:

```
location /stream/ {  
    proxy_buffering off;  
    proxy_pass http://backend_streamer;  
}
```

## Logging: custom formats and conditional logs

```
log_format main '$remote_addr - $remote_user [$time_local] "$request" '
                 '$status $body_bytes_sent "$http_referer" '
                 '"$http_user_agent" "$http_x_forwarded_for"';
access_log /var/log/nginx/access.log main;
```

Conditional logging (skip health checks):

```
map $request_uri $loggable {
    ~^/health$ 0;
    default 1;
}
access_log /var/log/nginx/access.log main if=$loggable;
```

## Error pages and maintenance mode

```
error_page 500 502 503 504 /50x.html;
location = /50x.html {
    root /usr/share/nginx/html;
}

# Maintenance
server {
    listen 80;
    server_name example.com;

    if (-f /var/www/maintenance.enable) {
        return 503;
    }

    error_page 503 @maintenance;
    location @maintenance {
        root /var/www/maintenance;
        try_files /index.html =503;
    }
}
```

## Split traffic with map (A/B testing / feature flag)

```
map $http_cookie $variant {
    ""      0;
    "ab=1"  1;
}

server {
    location / {
        if ($variant = 1) {
            proxy_pass http://blue_backend;
        }
        proxy_pass http://green_backend;
    }
}
```

Notes: map is efficient for runtime decisions.

## Sticky sessions (cookie based)

- Open-source: cookie-based sticky requires third-party module (nginx-sticky-module-ng) or use the `hash` directive/methods.
- NGINX Plus: has `sticky` directive for stable cookie-based persistence.

Example using `ip_hash` (simpler):

```
upstream backend {  
    ip_hash;  
    server 10.0.0.11:8000;  
    server 10.0.0.12:8000;  
}
```

## Using `include` for modular configs (recommended layout)

- `/etc/nginx/nginx.conf` includes `/etc/nginx/conf.d/*.conf` and `/etc/nginx/sites-enabled/*`.
- Put site-specific server blocks under `sites-available` and symlink into `sites-enabled`.
- Use `include snippets/` to store reusable pieces (proxy-headers, gzip-params, etc.).

## Lua example (embedding business logic)

- Requires `ngx_http_lua_module` (OpenResty or compiled module).

Simple snippet:

```
location /lua/ {
    content_by_lua_block {
        ngx.say("Hello from Lua: ", ngx.var.remote_addr)
    }
}
```

Notes: powerful but use with care (complexity and security considerations).

## Advanced: rate-limiting by key (user-based)

```
limit_req_zone $binary_remote_addr zone=perip:10m rate=5r/s;
limit_req_zone $cookie_userid zone=peruser:10m rate=1r/s;

server {
    location /api/ {
        limit_req zone=peruser burst=5 nodelay if=$cookie_userid;
        limit_req zone=perip burst=20;
        proxy_pass http://backend;
    }
}
```

## Troubleshooting tips

- Always run `nginx -t` after edits.
- Inspect logs: `/var/log/nginx/error.log` and `/var/log/nginx/access.log`.
- If config changes do not take effect, remember to `nginx -s reload` or `systemctl reload nginx`.
- Use `curl -I -H "Host: example.com" http://127.0.0.1` to test virtual hosts locally.

## **When to prefer NGINX Plus or modules**

- Active upstream health checks, advanced monitoring, dynamic reconfiguration → NGINX Plus
- Cookie-based sticky sessions without third-party modules → NGINX Plus
- Otherwise, open-source + third-party modules (or external tools) suffice for many setups.

## **Further reading**

- Official docs: <https://nginx.org/en/docs/>
- NGINX Plus features: <https://www.nginx.com/products/nginx/>
- OpenResty (nginx + lua): <https://openresty.org/>